

- Building Infrastructure
- Optimizing ML algorithm
- Integration

LINEAR REGRESSION WITH NUMPY & PYTHON

Build → Univariate Linear Regression.

TASK 142 Import Libraries and Load Data

GGPLOT → Graph Gallery

Univariate ↔ Linear Regression with one variable.

Data → Bike sharing data to decide which city to expand in.

Load Data.

Invest
3me.
data = pd.read_csv('name')
→ data.head()

data.info ← To get more insight.

TASK-3 VISUALISE THE DATA

We will use scatter plot as only 2 variables.

years for 'to predict'

ax = sns.scatterplot(x="Population", y="Profit")

data = data

ax.set_title("Profit in \$10000 vs City Population")

We will find a line of best fit. (Linear Regression)

TASK-4 Compute Cost J(θ)

Our Linear Regression powered by Gradient Descent.

Difference b/w Target Predicted

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

Cost/Error normalisation = 1 of data. Estimate Profit on x training Example Profit

2 is mathematical convenience & will cancel out when we compute.
m is length of data

h_{θ} is the hypothesis.

$$h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x,$$

Implementing Cost function

def cost-function (X, y, theta):

m = len(y)

y_pred = X.dot(theta)

errors = (y_pred - y)**2

return 1/(2*m) * np.sum(errors)

dot of x & theta

Dot product of θ and x .

m = data . Population . values . size.

Total data k size

x can be population but we need to add another dimension to accommodate the θ_0 intercept term. So, θ_0 is another feature of x .

$x = np.append(np.ones((m, 1)), data . Population . values . reshape((m, 1), axis = 1))$

first row of all 1

redefine the data with m rows k 1 column

changing to columns

$y = data . Profit . values . reshape((m, 1))$

theta = np.zeros((2, 1))

All zeroes.

rows columns

We added an extra dimension to our input or feature matrix x to accommodate the intercept term Explanation and set it to all ones **QUIZ**

$h_0(x)$ = predicted value of y at x **QUESTION**

$$y = mx + c \rightarrow \text{Eq. of line}$$

$$h_0(x) = \theta_1 x + \theta_0 \leftarrow \text{Intercept term i.e. where does line intersect y-axis}$$

So, I need to find values of θ_1 & θ_0

That's why my initial theta is all zero to begin with. And my theta only have 2 rows & 1 column i.e. 2 boxes. One represents θ_0 and other θ_1 .

Case I

normal $X \cdot \theta$

$$\theta_0 \begin{bmatrix} 0 \\ 0 \end{bmatrix}$$

Case II

appended $X \cdot \theta$

Case I

$$x \cdot \theta_0 + x \cdot \theta_1$$

Wrong equation

Case II

$$x_0 \theta_0 + x_1 \theta_1$$

we have

appended all 1

actual population values,

Overall Eqn $\rightarrow \theta_0 + x_1 \theta_1$, which is exactly what we want and that's why we appended and changed

BASIC IMAGE CLASSIFICATION (TENSORFLOW)

TASK-5 GRADIENT DESCENT

Reducing the error / cost function

Formula Used \rightarrow Batch Gradient Descent

def gradient_descent(x, y, theta, ^{learning rate}alpha, iterations):

m = len(y)

costs = []

to track all cost history
to see how error
changed.

+ ki dimension $\rightarrow (m, 2)$
use transpose $\rightarrow (2, m)$

for i in range(iterations):

is ki dimension
has $(m, 1)$

multiply $(2, m) \times (m, 1)$
result $\rightarrow (2, 1)$
to hamara theta

y_pred = x.dot(theta)

errors = np.dot(x.transpose(), (y_pred - y))

theta = theta - alpha * $\frac{1}{m}$ * errors

costs.append(cost_function(x, y, theta))

return theta, costs

theta, costs = gradient_descent(x, y, theta, alpha=0.01,
iterations=2000)

print("h(x) = {0} + {0}x".format(round(theta[0, 0], 2),
round(theta[1, 0], 2)))
intercept term

TASK-6 Visualize the cost function

from mpl_toolkits.mplot3d import Axes3D

theta_0 = np.linspace(-10, 10, 100)

theta_1 = np.linspace(-1, 4, 100)

creates a list of uniform numbers from to How many

cost_values = np.zeros((len(theta_0), len(theta_1)))

for i in range(len(theta_0)):
for j in range(len(theta_1)):

t = np.array([theta_0[i], theta_1[j]])

cost_values[i, j] = cost_function(x, y, theta)

fig = plt.figure(figsize=(12, 8)) *Size of graph.*

ax = fig.gca(projection='3d') *3d to 2D plotting*

surf = ax.plot_surface(theta_0, theta_1, cost_values)

cmap = 'viridis'

colour map *diversions*
colours
map.

fig.colorbar(surf, shrink=0.5, aspect=5)

plt.xlabel("\$\theta_0\$")

plt.ylabel("\$\theta_1\$")

ax.set_zlabel("\$J(\theta)\$") *cost function*

ax.view_init(30, 330) *rotating for better view*

TASK-7 PLOTTING THE CONVERGENCE

Task \rightarrow Plot the costs found above.

```
plt.plot(costs)
plt.xlabel("Iterations")
plt.ylabel("$J(\theta)$")
plt.title("Values of the cost function over  
iteration of gradient descent")
```

TASK-8 Training Data with Linear Regression Fit.

theta.shape $\rightarrow (2, 1)$

theta \rightarrow array $\begin{bmatrix} [-3.7], [1.182] \end{bmatrix}$

theta = np.squeeze(theta) \rightarrow Remove extra dim.

Range of values.

x-value = $[x \text{ for } x \text{ in range}(5, 25)]$

y-value = $[(x * \text{theta}[1] + \text{theta}[0]) \text{ for } x \text{ in } x\text{-value}]$

The equation i.e

$$y = h_{\theta}(x) = \theta^T x = \theta_0 + \theta_1 x$$

loop for all values of x

Important
if is a problem
we mostly forget
But Hard to Debug

Display
Line

```
sns.lineplot(x-value, y-value)  
plt.xlabel("name")  
plt.ylabel("name")
```

TASK-9 Inference using optimised θ value
or using line to make predictions.

```
def predict(x, theta)
```

```
    y_pred = np.dot(theta.transpose(), x)
```

```
    return y_pred
```

Give first value of y to predict profit. ^{because θ_0 was multiplied by 1} Actual x .

```
y_pred - 1 = predict(np.array([1, 4]) theta) * 10000
```

```
print("The estimated profit is " y_pred - 1)
```

```
y_pred - 2 = predict(np.array([1, 8.3]) theta) * 10000
```